



# New Relic Unix Monitor

Install and Use Guide

System-Level Monitoring for AIX, Linux, OSX/macOS & Solaris/SunOS

## Disclaimer

New Relic has this integration to enable monitoring of this technology. This integration is provided AS-IS WITHOUT WARRANTY OR SUPPORT, although you can report issues and contribute to this integration via GitHub.

## Document Change History

Version	Date	Change
0.2.5	2022-Mar-08	Added fedramp flag
1.0.1	2023-Jan-13	Added release notes

## Table of Contents

<b>Requirements</b>	<b>4</b>
Supported Operating Systems	4
<b>Installation &amp; Usage Overview</b>	<b>5</b>
<b>Configuration</b>	<b>6</b>
<b>Plugin configuration</b>	<b>6</b>
Global settings	6
Agent settings	7
Sample Plugin.json	8
<b>Command Configuration</b>	<b>9</b>
eventType	9
command	9
checkAllRegex	9
lineLimit	10
interval	10
mappings	11
Example command configuration	12
<b>Startup script</b>	<b>13</b>
Usage	13
<b>Advanced configurations</b>	<b>14</b>
Deploying Dashboards from separate server/desktop	14
Fix for using the WebSphere JDK	14
Fix for using Solaris 10	14
Guided script driver	15
Configuration	15
Sample Regular Expression Java Program	16
Release Notes	18

## Requirements

- A New Relic [account](#)
- Supported Unix server to be monitor
- Java JRE/JDK v1.6 or later
- Network access to New Relic (proxies are supported, see details below)
- For Dashboard Installation: *curl* or *wget* installed

## Supported Operating Systems

- o AIX 7.x
- o Linux - All sorts, including on ARM processors (such as Raspberry Pi) and z/Linux
- o OSX / MacOS 10.9 ('Mavericks') and above
- o Solaris/SunOS 10.x and 11.x

## Installation & Usage Overview

1. Download the latest version of the agent.
2. Gunzip & untar on Unix server that you want to monitor
3. Configure and customize the following.
  - a. [Plugin configuration](#) *\*required*  
*plugin.json to set account ID, keys and static attributes*
  - b. [Command configuration](#) *(optional)*  
*plugin-commands-\*.json files to configure additional Operating system commands*
  - c. [Startup script](#) *(optional)*  
*pluginctl.sh to set java and monitor path*
4. Run “**./pluginctl.sh start**” from command-line to start the monitor
5. Check logs (in logs directory by default) for errors
6. Login to New Relic UI and find your data in Insights.
  - a. In the data explorer, look for custom event types that start with "unixMonitor:"
  - b. Possible event types (for out-of-the-box commands): unixMonitor:Disk, unixMonitor:DiskIO, unixMonitor:NetworkIO, unixMonitor:Process, unixMonitor:Stats, unixMonitor:Vmstat

## Configuration

### Plugin configuration

Plugin configuration is governed by the ***“plugin.json”*** file. This file is used to set global and agent settings for the monitor. A full example of the possible fields in *plugin.json* can also be found in *plugin-full-example.json*

### Global settings

#### ***“global”***

- OS (default: auto):  
Used to determine which commands to run and how to parse them. Leave set to auto to have the plugin figure that out (which normally works).
- account\_id:  
New Relic account ID - the 6- or 7- digit number in the URL when you're logged into the account of your choosing.
- fedramp:  
A true or false string to indicate that the target is the New Relic Fedramp-authorized endpoint.

#### ***“infra\_mode”***

This object is to be defined when we use this monitor to push the data using new relic infrastructure agent

- "rpc\_listener\_port"  
RPC listener port on which is used to push the data infrastructure agent.

#### ***“insights\_mode”***

This object needs to be defined when we are directly posting to “Insights”

- insights\_insert\_key  
You must create an [Insights Insert key, as described here](#).

### "dashboards"<sup>1</sup>

This object needs to be defined if default dashboards need to be deployed using the `pluginctl.sh` script

- `admin_api_key`  
specify the [Admin API key, as described here.](#)
- `integration_guid`  
Default is `UNIX.Infra.Monitor`.
- `dashboard_install`  
Default is `command_line`.

**Note:** DO NOT DELETE OR CHANGE `integration_guid` AND `dashboard_install` UNLESS OTHERWISE INSTRUCTED. Both are required but must be left to their default values.

### "proxy"

If using a proxy, the optional proxy object should be added to the global object in `plugin.json`.

- `proxy_host`  
Hostname for proxy
- `proxy_port`  
Proxy port
- `proxy_username`  
Proxy username
- `proxy_password`  
Proxy password

## Agent settings

### "agents"

- `name`  
If set to `auto`, the plugin will use that server's hostname. Otherwise, sets the hostname and `agentName` to whatever is set here.
- `static` (optional)  
An object containing static attributes (as name-value pairs) you want to appear in every event from this plugin

---

<sup>1</sup> \*\* requires `curl` or `wget` installed

## Sample Plugin.json

```
{
  "global": {
    "OS": "auto",
    "account_id": "enter_NR_account_ID",
    "fedramp": "false",
    "insights_mode": {
      "insights_insert_key": "enter_insights_insert_key"
    },
    "dashboards": {
      "admin_api_key": "enter_admin_api_key",
      "integration_guid": "UNIX.Infra.Monitor",
      "dashboard_install": "command_line"
    },
    "proxy": {
      "proxy_host": "enter_proxy_host",
      "proxy_port": 5443,
      "proxy_username": "enter_proxy_username",
      "proxy_password": "enter_proxy_password"
    }
  },
  "agents": [
    {
      "name": "auto",
      "static": {
        "attribute1": "attribute1_value",
        "attribute2": 12345
      }
    }
  ]
}
```



## Command Configuration

Command configuration (plugin-commands-\*.json) is used to configure the various commands that are used to collect the metrics from the target operating system. The kit contains default files for each supported Operating System (OS) configured to pull out the basic metrics. These files can be extended to pull additional metrics for the operating system by configuring additional commands in these files.

The files consist of the following parameters that can be configured in json format for a command in OS specific.

### eventType

This attribute defines the type of event under which the metrics will be listed. The value listed here will be prepended with "unixMonitor:" and will be listed in data explorer under custom event types. In the [example configuration](#), the eventType of **"File"** is listed as **unixMonitor:File** under custom event types in Insights UI data explorer

### command

This attribute defines the actual command that needs to be run to collect a metric. In [example configuration](#) the following command is run

**"ls -l /opt/New\_Relic/newrelic-unix-monitor/config/plugin.json"**

The output from this command is parsed by the [expression](#) defined later in [mapping](#) section

```
$ ls -l /opt/New_Relic/newrelic-unix-monitor/config/plugin.json
-rw-r--r-- 1 user staff 355 Dec 19 07:22
/opt/New_Relic/newrelic-unix-monitor/config/plugin.json
```

### checkAllRegex

This attribute can have two values (true/false). This governs if all the regular expressions listed under mappings will be matched. This is only for special cases where we need to iterate all the expressions. In [example configuration](#) this is set to false

## lineLimit

This attribute defines how many lines need to be parsed. Default "0" will parse all the lines This is useful when we need to parse only a portion of command output for regular expression matching. In [example configuration](#) this is set to 0

## interval

This attribute defines in minutes the interval between running the configured command. By default, interval 0, the commands would be run each time a harvest cycle is run. In [example configuration](#) this is set to 1440 i.e. 1day.

## mappings

The mapping object defines how the output of the command needs to be read and what metrics they need to be matched to. Multiple mapping can be defined for a command. This object has sub attributes defined below

### expression

This attribute defines the regular expression that is used to parse and match the line output. This expression needs to define the capturing group '(' to map them to metrics defined below. For more information refer java regular expression pattern matching <sup>2</sup>. In the [example configuration](#) below the following expression is used

```
(\\S+)\\s+(\\d+)\\s+(\\S+)\\s+(\\S+)\\s+(\\d+)\\s+(\\S+\\s+\\d+\\s+\\d+:\\d+)\\s+(\\S+)
```

This expression breaks down the output of the command into groups<sup>3</sup> as shown below in table

Match #	Group index	Start index	End index	Group content
1	0	0	112	-rw-r--r-- 1 user staff 355 Dec 19 07:22 /opt/New_Relic/newrelic-unix-monitor/config/plugin.json
1	1	0	10	-rw-r--r--
1	2	12	13	1
1	3	14	20	user
1	4	22	27	staff
1	5	29	32	355
1	6	33	45	Dec 19 07:22
1	7	46	112	/opt/New_Relic/newrelic-unix-monitor/config/plugin.json

<sup>2</sup> Java regex pattern matching @ <https://docs.oracle.com/javase/7/docs/api/java/util/regex/Pattern.html>

<sup>3</sup> To parse the output for configuration use online [java regex testers](#)

## metrics

This object defines the metrics that are parsed using the regular expression and are grouped. The metrics must be defined in sequential manner i.e. group 1 followed by group 2 and so on. The sub objects must define the two parameters listed below.

### name

This attribute defines the name of the metric. In the example below for the first group name is listed as "**File.permissions**" the value that will map to this metric from the regular expression is "**-rw-r--r--**"

### type

This attribute defines the type of the metric. In most cases this will be listed as "**NORMAL**" although for complex cases "DELTA" is defined that reports the delta between the current and previous run of harvest cycle.

## Example command configuration

```
{
  "eventType": "File",
  "command": "ls -l
/opt/New_Relic/newrelic-unix-monitor/config/plugin.json",
  "checkAllRegex": false,
  "lineLimit": 0,
  "interval": 1440,
  "mappings": [
    {
      "expression":
"((\\S+)\\s+(\\d+)\\s+(\\S+)\\s+(\\S+)\\s+(\\d+)\\s+(\\S+\\s+\\d+\\s+\\d+:\\d+)\\s+(\\S+)",
      "metrics": [
        {
          "name": "File.permissions",
          "type": "NORMAL"
        },
        {
          "name": "File.links",
          "type": "NORMAL"
        },
        {
          "name": "File.owner",
          "type": "NORMAL"
        },
        {
          "name": "File.grouup",
          "type": "NORMAL"
        },
        {
          "name": "File.size",
          "type": "NORMAL"
        },
        {
          "name": "File.Date",
          "type": "NORMAL"
        },
        {
          "name": "File.Path",
          "type": "NORMAL"
        }
      ]
    }
  ]
}
```

Note : A [Sample Java code](#) with multiple sample expressions is given below to help defining the correct regular expression for your output.

## Startup script

The startup script “pluginctl.sh” is used to control the execution of the monitor, This script can also be modified to set various environmental parameters. The two most widely used configurations are

- o PLUGIN\_JAVA to set the location of Java on your server (including the "java" filename)
- o PLUGIN\_PATH to set the location of the Unix monitor

## Usage

The startup script provides the following options.

```
Usage: ./pluginctl.sh [status|start|stop|stopremlogs|restart|dashboards]
```

### **status**

Query the status of the monitor.

### **start**

Start the monitor

### **stop**

stop the monitor

### **stopremlogs**

Stop the monitor and remove logs

### **restart**

Restart the monitor

### **dashboards**

Deploy [dashboards](#). For more details check advanced configuration

## Advanced configurations

### Deploying Dashboards from separate server/desktop

You can initiate the dashboard install from a standalone machine (i.e. a tools server or your own mac, linux or cygwin laptop/desktop), you will need the following:

- pluginctl.sh
- config/plugin.json (including path) with the dashboard object filled

To install, run `./pluginctl.sh dashboards`.

### Fix for using the WebSphere JDK

If you are using the JDK that is packaged with WebSphere and see an exception in the logs like below, it is due to attempting to use the WebSphere SSL Factory instead of the IBM JSSE packages.

```
ERROR com.newrelic.metrics.publish.binding.Request - An error occurred
communicating with the New Relic service
java.net.SocketException: java.lang.ClassNotFoundException: Cannot find
the specified class
com.ibm.websphere.ssl.protocol.SSLSocketFactory
```

If so, uncomment the following line in `pluginctl.sh` and restart the plugin.

```
# USE_IBM_JSSE=true
```

### Fix for using Solaris 10

If you see the following error, it may be because the Bourne shell does not support certain syntax in the installer script.

```
pluginctl.sh: syntax error at line 240: `admin_api_key=$' unexpected
```

If so, use the Korn shell or bash (if available). Both were tested successfully in Solaris 10.

## Guided script driver

A new guided script driver has been added in the *r1.0.1* release that will run the configured scripts periodically and send the script's status to New Relic .

### Configuration

**Step 1 :** From the folder `newrelic-unix-monitor` run the following command

```
# ./gen_scriptdriver_script.sh

*****ScriptDriver Options*****

*. Press 1 to update the scripts' folder location:
    [/export/home/gulab/data/newrelic-unix-monitor/scripts]
*. Press 2 to continue
*. Press 3 to exit

Enter your choice:
```

**Step 2 :** Define the folder location where all your scripts are available and those are required to be run on a periodic basis . Press 1 , if you need to update the folder location

**Step 3 :**Once the folder location has been selected then Press 2. This will generate the driver script `script_driver.sh` under `newrelic-unix-monitor`.

```
.
.
Info: Driver file created. Please create a cron job as per your required frequency

Here is an example to run every minute

* * * * * /export/home/gulab/data/newrelic-unix-monitor/script_driver.sh
          >/dev/null 2>&1

You may exit !!

*****ScriptDriver Options*****

*. Press 1 to update the scripts' folder location:
    [/export/home/gulab/data/newrelic-unix-monitor/scripts]
*. Press 2 to continue
*. Press 3 to exit

Enter your choice:
```

**Step 4 :**Press 3 and exit. Give execute permissions to the script `script_driver.sh`

**Step 5.** Now update the cron job as per the instructions from the output above.

run `crontab -e` and add the lines from the output.

```
#  
# Copyright (c) 1989, 2016, Oracle and/or its affiliates. All rights reserved.  
#  
# The root crontab should be used to perform accounting data collection.  
#  
#  
10 3 * * * /usr/sbin/logadm  
30 3 * * * [ -x /usr/lib/gss/gsscred_clean ] && /usr/lib/gss/gsscred_clean  
* * * * * /export/home/gulab/data/newrelic-unix-monitor/script_driver.sh  
                >/dev/null 2>&1
```

**Step 6.** Restart Cron Service. This will take the new job and start running as per the configuration.

```
svcadm restart cron
```

**Note:** All the scripts' status will be collected in `/tmp/script_status.out` and the unix monitor will read this file during each harvest cycle and remove this file.

- New Metrics Group **`unixMonitor:UnixScriptStats`** has been added as part of Solaris 11 configuration.
- Below query may be used to fetch the data:
  - `Select instance, stauscode, executedat FROM `unixMonitor:UnixScriptStats``

## Sample Regular Expression Java Program

A Sample [PatternMatcher.java](#) program has been added to help you to understand how to parse a line. There are multiple sample lines with regular expressions are given in the program and you may uncomment and run the same. Once the parsing is done correctly then the same working regular expression may be used in newrelic-unix-monitor's json file.



## Release Notes

**Date: 13-Jan-2023**

**Version: 1.0.1**

The release 1.0.1 is a major release with multiple fixes and enhancements for Solaris 11.x platform, and this release has been certified on Oracle Solaris 11.4.

Here are the [key enhancements and fixes](#):

- `cpu.iowait` has been updated with `cpu.stolen` in the config file based on the command output. The `cpu.iowait` is no more available in command output
  - Below query may be used to fetch the data:
    - `Select cpu.idle,cpu.kernel,cpu.stolen,cpu.swap,cpu.user from `unixMonitor:Stats``
- This release has been enhanced to supports multiple commands separated by `| && ;`
  - Example nested command like `echo ::memstat | mdb -k` is now supported.
  - New Metrics Group ``unixMonitor:Memstat`` has been added as part of Solaris 11 configuration to collect the memory stats.
- Logical Domains Manager (`ldm ls -p`) command output is now being sent to New Relic.
  - Refer `newrelic-unix-monitor/scripts/getLdmMemStat.sh` and follow the instructions.
  - New Metrics Group ``unixMonitor:LDMMemstat`` has been added as part of Solaris 11 configuration.
- Unix Faults (`fmadm list -s`) command output is now being sent to New Relic.
  - Refer `newrelic-unix-monitor/scripts/getUnixFaults.sh` and follow the instructions.
  - New Metrics Group ``unixMonitor:UnixFaults`` has been added as part of Solaris 11 configuration.
- All the services status (`svcs -av`) command output is now being sent to New Relic.
  - New Metrics Group ``unixMonitor:ServiceStatus`` has been added as part of Solaris 11 configuration.
- For the disk statistics the mount point is now captured as part of the *instance* field.
  - Below query may be used to fetch the data:
    - `Select instance FROM `unixMonitor:Disk``

- Existing Metrics Group **unixMonitor:DiskIO** has been fixed to capture the required Disk IO stats.
- Attribute values are now trimmed towards spaces to allow queries to support where clauses..
  - Below query may be used to fetch the data:
    - `Select * from `unixMonitor:Memstat` where instance='ZFS'`
- A new [guided script driver](#) has been added that will run the scripts periodically and send the script's status to New Relic
  - New Metrics Group **`unixMonitor:UnixScriptStats`** has been added as part of Solaris 11 configuration.
  - Below query may be used to fetch the data:
    - `Select instance, stauscode, executedat FROM `unixMonitor:UnixScriptStats``
- [A Sample Java program](#) has been provided with multiple samples of line output and regular expressions to define and test your regular expressions for a given system output .
  - This will help to define a new New Metrics Group and send the data to the New Relic platform.
- [Solaris Log Integration](#) with New Relic has been introduced .
  - A [Step by Step guide](#) has been provided as part of this release.
- Process statistics has been [enhanced](#) to collect the user, time and nlwp
  - Updated Metrics Group **`unixMonitor:Process`** to capture user, time and nlwp
  - This will help to Get the Hung Process running on Solaris
- [Fixed](#) Metrics Group **unixMonitor:fileWatcher** Example configuration to capture all fields correctly .